

Supervised Spam Classification

“Comparing Effectivity and Robustness of Sequential and Non-Sequential Machine-Learning Models”

Atreya Shankar, Cognitive Systems (M.Sc.)

BM2: Intelligent Data Analysis (IDA)

University of Potsdam, SoSe 2019

Prof. Dr. Tobias Scheffer

August 08. 2019

Table of Contents

1 Introduction

2 Methodologies

3 Results

4 Evaluation

5 Conclusions

6 Bibliography

Introduction

Intelligent Data Analysis

Exam: Spam (Project 5)

Summer term 2018

Prof. Tobias Scheffer
Dr. Paul Prasse
Silvia Makowski
Dr. Lena Jäger

This project is part of the exam *Intelligent Data Analysis*. Each project assignment is to be resolved by a single student on his/her own. The student is supposed to present the solution as part of the oral exam. The student is required to present a printed version of the Python code together with diagrams, tables, etc. that summarize the results. The specific way of how the project is presented is up to the student's choice.

Problem setting

You have been hired by the IT department of a medium-sized company to train an email spam filter which should mark the incoming emails of all employees as spam or non-spam. The emails are parsed by a module and converted into the bag-of-words representation. A total of 57,173 different words (features) are distinguished. The aim of the filter is to identify a maximum number of spam emails, with a maximum of 0.2% of all legitimate emails being classified incorrectly. In addition, the company wants to make a statement about the effectiveness of the filter on future emails, i.e., what percentage of incoming spam emails will be identified in the future.

Aufgabe

From the employees' inboxes, 10,000 emails were extracted as training data (see emails.mat). Let X be the training data with the associated class labels Y (+1 stands for *spam*, -1 means *non-spam*). Identify a suitable learning technique for constructing a spam filter and implement it in Python. Train and evaluate the model. Make a statement about the expected quality of the filter and make sure that no more than 0.2% of all legitimate emails are filtered. For this purpose, plot a precision/recall curve and mark the model's position on the curve (for the selected threshold). Briefly motivate and document all the steps you have taken.

Fig. 1. Spam project description

- Project description proposes using data in “emails.mat” file with 10k instances and ~50k features
- Bag-of-words form of data, which would only work for non-sequential learning
- Enron-spam pre-processed text data derived from Enron Corporation scandal; subset of employees' emails became publicly available (Metsis, Androutsopoulos, and Paliouras, 2006)
- Consists of 33,716 text-based emails; 16,545 “ham” and 17,171 spam instances

Objectives

- Utilize enron-spam emails database to implement both sequential and non-sequential supervised classifiers
- Meet project requirement to develop a classifier that attains 99.8% recall on “ham” emails
- Provide input into recall values for future spam emails given selected optimal threshold
- Additionally, provide insights into effectivity and robustness of sequential and non-sequential models

Overview

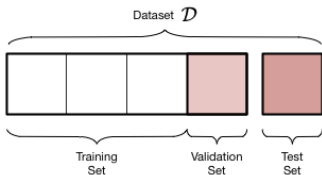


Fig. 2. Data splitting schematic (Ziganto, 2018)

- Non-sequential model: Support Vector Machine (SVM)
- Sequential model: CNN-LSTM with word/character embeddings
- Due to time limitations, K-fold cross-validation was omitted
- Compromise: train/validate/test on the same subsets of data for fair comparison
- $(\text{Train} \cup \text{Validation}):\text{Test} \implies 70:30$
- $\text{Train}:\text{Validation} \implies 85:15$

Non-Sequential Model: Support Vector Machine (SVM)

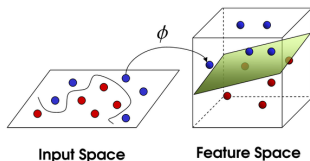


Fig. 3. Support Vector Machine (SVM) schematic (Joshi, 2012)

- Pre-processing text to normalized bag-of-words representation with $|V| = 5,000$ words
- Sklearn's `SGDClassifier` with Mini-Batch SGD and early stopping
- Linear and approximated RBF Kernel (`RBFSampler`)
- Grid-search over batch-size, regularization term α , RBF kernel γ , and number of sampling components for `RBFSampler`

Sequential Model: CNN-LSTM (Words)




Fig. 4. Keras schematic for CNN-LSTM (Words)

- Pre-processing text to padded/clipped integer encoded tokens with $|V| = 5,000$ words
- 1-dimensional CNN with varying filters to enrich sequential features; LSTM cell to capture short and long-term sequential relationships; dropout regularization for model robustness
- Grid-search over embedding dimensions, dropout rate, batch-size and learning rate
- Learning both with and without pre-trained GloVe word vectors (~ 6 billion tokens)

Sequential Model: CNN-LSTM (Words+Characters)

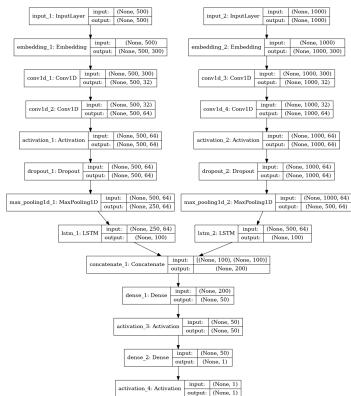


Fig. 5. Keras schematic for CNN-LSTM (Words+Characters)

- Using character sequences to overcome unknown token issue; same general architecture as before
- Grid-search over embedding dimensions, dropout rate, batch-size and learning rate
- Learning both with and without pre-trained GloVe word vectors (~6 billion tokens)
- Approximating GloVe character embeddings by averaging over character-containing word vectors (Woolf, 2017)

Grid-search optimal models

Classifier	Training F_1	Validation F_1	Test F_1
SVM (Linear Kernel)	0.9880	0.9820	0.9836
SVM (Approximated RBF Kernel)	0.4353	0.3494	0.3437
CNN-LSTM (Words)	0.9904	0.9791	0.9753
CNN-LSTM (Words+Characters)	0.9934	0.9841	0.9808
CNN-LSTM (Words+GloVe)	0.9998	0.9903	0.9902
CNN-LSTM (Words+Characters+GloVe)	0.9992	0.9909	0.9902

Table 1: Summary of grid-search optimal models; zero rule classifier baseline is 50.9%

- Both sequential and non-sequential models achieve high F_1 test scores

“Ham” Relative Importance Analysis (SVM)

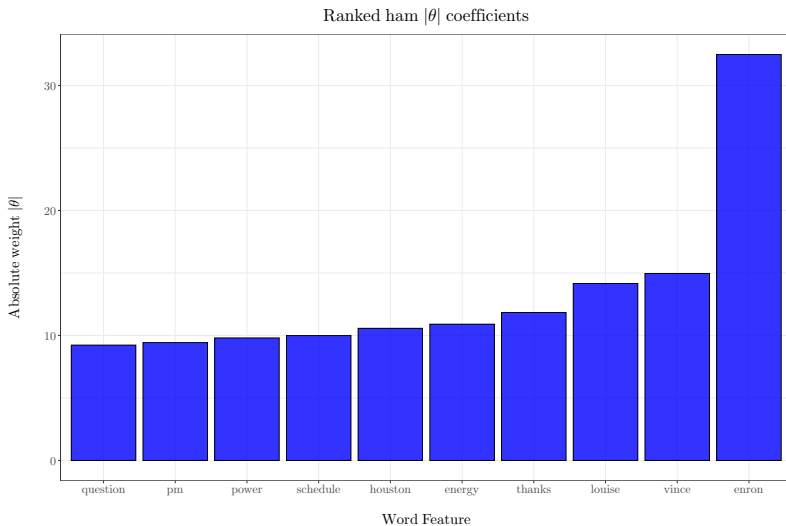


Fig. 6. Relative importance analysis for SVM (linear kernel)

Spam Relative Importance Analysis (SVM)




Fig. 7. Relative importance analysis for SVM (linear kernel)

Optimal Threshold Analysis

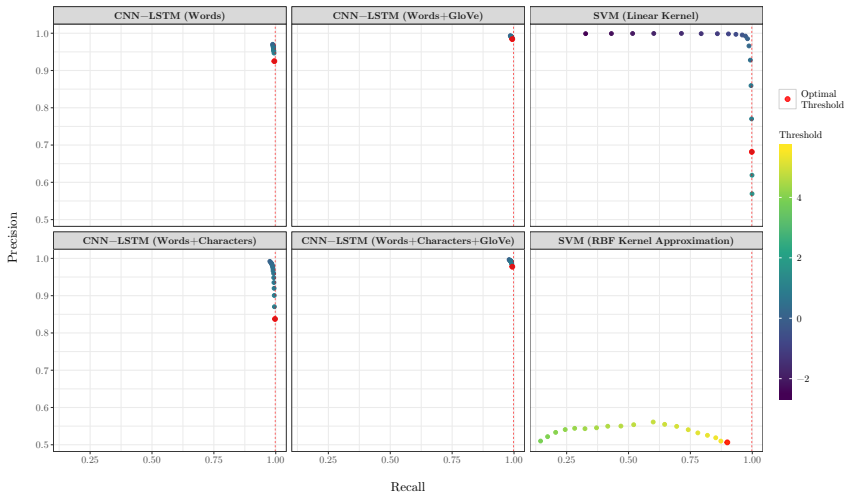


Fig. 8. Precision-Recall curve (ham label) for optimal threshold analysis

Optimal Threshold Performance

Classifier	Threshold	Recall (Spam)	Recall (Ham)
SVM (Linear Kernel)	1.171	0.5509	0.9982
SVM (Approximated RBF Kernel)	5.552	0.1558	0.8991
CNN-LSTM (Words)	0.9444	0.9224	0.9943
CNN-LSTM (Words+Characters)	0.9444	0.8137	0.9971
CNN-LSTM (Words+GloVe)	0.9444	0.9846	0.9929
CNN-LSTM (Words+Characters+GloVe)	0.9444	0.9781	0.9930

Table 2: Results of optimal threshold analysis

- Trade off between ham and spam recall; either accept low recall for spam or compromise with lower recall for ham

Blind Dataset Performance (SMS Spam)

Classifier	Blind F_1
SVM (Linear Kernel)	0.4688
SVM (Approximated RBF Kernel)	0.1785
CNN-LSTM (Words)	0.5090
CNN-LSTM (Words+Characters)	0.4416
CNN-LSTM (Words+GloVe)	0.2913
CNN-LSTM (Words+Characters+GloVe)	0.3017

Table 3: Results of blind data test; zero rule classifier obtains 87% due to class imbalance

- Words-based models perform best on blind dataset (albeit worse than zero rule classifier)
- Considering sequential nature of data contributes to robustness

Conclusions

- For spam detection, both sequential and non-sequential models are effective
- Trade-off exists between “ham” and spam recall; an informed decision must be made
- Both words-based sequential and non-sequential modes tend to be robust to new datasets; although sequential models tend to carry richer and more discriminating features
- High cost of training CNN-LSTM; perhaps not economical for a company to deploy a GPU on IMAP server
- **SVM would be a more efficient and scalable option**

Improvements to Embeddings in CNN-LSTM

- Separate pipeline for character sequences leads to symbolic overfitting on types of datasets
- Can overcome unknown tokens but contributes uncertainty in terms of dialects and expressions
- Zhao (2018) proposes a bidirectional LSTM to enrich word vector features
- This could address the unknown token issue without leading to overfitting on entire character sequences




Fig. 9. Improved word-character embedding model (Zhao, 2018)

Bibliography I

Joshi, Prateek (2012). URL:

<https://prateekvjoshi.com/2012/08/24/support-vector-machines/>.

Metsis, Vangelis, Ion Androutsopoulos, and Georgios Paliouras (2006). "Spam filtering with naive bayes-which naive bayes?" In: *CEAS*. Vol. 17. Mountain View, CA, pp. 28–69.

Woolf, Max (2017). *char-embeddings (Github)*. URL:

<https://github.com/minimaxir/char-embeddings>.

Zhao, HG (2018). *keras-word-char-embs (Github)*. URL:

<https://github.com/CyberZHG/keras-word-char-embd>.

Bibliography II

Ziganto, David (2018). *Model Tuning (Part 2 - Validation & Cross-Validation)*. URL:

<https://dziganto.github.io/cross-validation/data%20science/machine%20learning/model%20tuning/python/Model-Tuning-with-Validation-and-Cross-Validation/>.